

## Project Title: Air Quality Monitoring and Pollution Prediction

Name: Kibon Kiprono Solomon

Date: 29/08/2023

**Problem Statement:** Air pollution is a growing concern in urban areas, with adverse effects on public health and the environment. Developing an air quality monitoring and pollution prediction system using machine learning can help raise awareness, inform policy decisions, and enable citizens to take protective measures.

**Project goal:** The aim of this project is to Create a machine learning system that monitors air quality in real-time, predicts pollution levels, and provides actionable insights for individuals and authorities to mitigate the impact of air pollution.

In [3]:

```
#import necessary Libraries
import numpy as np
import pandas as pd
import plotly.express as px
import seaborn as sns
from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV, RandomizedSearchCV
from sklearn.linear_model import LinearRegression, Ridge, Lasso
from sklearn.preprocessing import PolynomialFeatures, StandardScaler
from sklearn.svm import SVR
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from scipy.stats import randint
```

In [4]:

```
#Load the dataset
data=pd.read_csv('Air_quality_data.csv')
# Display the first few rows of the dataset
data.head(10)
```

Out[4]:

	City	Date	PM2.5	PM10	NO	NO2	NOx	NH3	CO	SO2	O3	Benze
0	Ahmedabad	2015-01-01	NaN	NaN	0.92	18.22	17.15	NaN	0.92	27.64	133.36	0.
1	Ahmedabad	2015-01-02	NaN	NaN	0.97	15.69	16.46	NaN	0.97	24.55	34.06	3.
2	Ahmedabad	2015-01-03	NaN	NaN	17.40	19.30	29.70	NaN	17.40	29.07	30.70	6.
3	Ahmedabad	2015-01-04	NaN	NaN	1.70	18.48	17.97	NaN	1.70	18.59	36.08	4.
4	Ahmedabad	2015-01-05	NaN	NaN	22.10	21.42	37.76	NaN	22.10	39.33	39.31	7.
5	Ahmedabad	2015-01-06	NaN	NaN	45.41	38.48	81.50	NaN	45.41	45.76	46.51	5.
6	Ahmedabad	2015-01-07	NaN	NaN	112.16	40.62	130.77	NaN	112.16	32.28	33.47	0.
7	Ahmedabad	2015-01-08	NaN	NaN	80.87	36.74	96.75	NaN	80.87	38.54	31.89	0.
8	Ahmedabad	2015-01-09	NaN	NaN	29.16	31.00	48.00	NaN	29.16	58.68	25.75	0.
9	Ahmedabad	2015-01-10	NaN	NaN	NaN	7.04	0.00	NaN	NaN	8.29	4.55	0.

In [5]:

```
# Check the distribution of the target variable (AQI_Bucket)
data['AQI_Bucket'].value_counts()
```

Out[5]:

```
Moderate      8829
Satisfactory  8224
Poor          2781
Very Poor     2337
Good          1341
Severe        1338
Name: AQI_Bucket, dtype: int64
```

In [6]:

```
#Remove null values from the target variable
data=data[data['AQI'].notna()]
data.head()
```

Out[6]:

	City	Date	PM2.5	PM10	NO	NO2	NOx	NH3	CO	SO2	O3	Benzene
28	Ahmedabad	2015-01-29	83.13	NaN	6.93	28.71	33.72	NaN	6.93	49.52	59.76	0.02
29	Ahmedabad	2015-01-30	79.84	NaN	13.85	28.68	41.08	NaN	13.85	48.49	97.07	0.02
30	Ahmedabad	2015-01-31	94.52	NaN	24.39	32.66	52.61	NaN	24.39	67.39	111.33	0.22
31	Ahmedabad	2015-02-01	135.99	NaN	43.48	42.08	84.57	NaN	43.48	75.23	102.70	0.42
32	Ahmedabad	2015-02-02	178.33	NaN	54.56	35.31	72.80	NaN	54.56	55.04	107.38	0.42

In [7]:

```
# Display data types and non-null counts of columns
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 24850 entries, 28 to 29530
Data columns (total 16 columns):
#   Column      Non-Null Count  Dtype
---  -
0   City         24850 non-null  object
1   Date         24850 non-null  object
2   PM2.5        24172 non-null  float64
3   PM10         17764 non-null  float64
4   NO           24463 non-null  float64
5   NO2          24459 non-null  float64
6   NOx          22993 non-null  float64
7   NH3          18314 non-null  float64
8   CO           24405 non-null  float64
9   SO2          24245 non-null  float64
10  O3           24043 non-null  float64
11  Benzene      21315 non-null  float64
12  Toluene      19024 non-null  float64
13  Xylene       9478 non-null   float64
14  AQI          24850 non-null  float64
15  AQI_Bucket   24850 non-null  object
dtypes: float64(13), object(3)
memory usage: 3.2+ MB
```

In [8]:

```
#sum of missing values in each column  
data.isna().sum()
```

Out[8]:

```
City          0  
Date          0  
PM2.5        678  
PM10         7086  
NO           387  
NO2          391  
NOx          1857  
NH3          6536  
CO           445  
SO2          605  
O3           807  
Benzene      3535  
Toluene      5826  
Xylene      15372  
AQI          0  
AQI_Bucket   0  
dtype: int64
```

In [9]:

```
# Calculate the percentage of missing values in each column  
missing_percentage = (data.isnull().sum() / len(data)) * 100  
missing_percentage
```

Out[9]:

```
City          0.000000  
Date          0.000000  
PM2.5        2.728370  
PM10         28.515091  
NO           1.557344  
NO2          1.573441  
NOx          7.472837  
NH3          26.301811  
CO           1.790744  
SO2          2.434608  
O3           3.247485  
Benzene      14.225352  
Toluene      23.444668  
Xylene      61.859155  
AQI          0.000000  
AQI_Bucket   0.000000  
dtype: float64
```

In [10]:

```
#remove columns with more than 50% missing values  
data.drop('Xylene',axis=1,inplace=True)
```

In [11]:

```
#fill the rest of missing values with median
data.fillna(data.median(),inplace=True)
data.head()
```

C:\Users\hp\AppData\Local\Temp\ipykernel\_15272\3294820205.py:2: FutureWarning: The default value of numeric\_only in DataFrame.median is deprecated. In a future version, it will default to False. In addition, specifying 'numeric\_only=None' is deprecated. Select only valid columns or specify the value of numeric\_only to silence this warning.

```
data.fillna(data.median(),inplace=True)
```

Out[11]:

	City	Date	PM2.5	PM10	NO	NO2	NOx	NH3	CO	SO2	O3	Benzer
28	Ahmedabad	2015-01-29	83.13	96.18	6.93	28.71	33.72	16.31	6.93	49.52	59.76	0.0
29	Ahmedabad	2015-01-30	79.84	96.18	13.85	28.68	41.08	16.31	13.85	48.49	97.07	0.0
30	Ahmedabad	2015-01-31	94.52	96.18	24.39	32.66	52.61	16.31	24.39	67.39	111.33	0.0
31	Ahmedabad	2015-02-01	135.99	96.18	43.48	42.08	84.57	16.31	43.48	75.23	102.70	0.0
32	Ahmedabad	2015-02-02	178.33	96.18	54.56	35.31	72.80	16.31	54.56	55.04	107.38	0.0

In [12]:

```
#descriptive statistics
data.describe()
```

Out[12]:

	PM2.5	PM10	NO	NO2	NOx	NH3
count	24850.000000	24850.000000	24850.000000	24850.000000	24850.000000	24850.000000
mean	66.966637	112.102860	17.502312	28.870163	31.645675	21.865639
std	62.283431	76.325808	22.266346	24.447523	29.629575	22.460343
min	0.040000	0.030000	0.030000	0.010000	0.000000	0.010000
25%	29.560000	71.780000	5.720000	12.090000	14.030000	11.280000
50%	48.785000	96.180000	9.910000	22.100000	23.680000	16.310000
75%	79.507500	122.957500	19.710000	37.910000	38.170000	24.710000
max	914.940000	917.080000	390.680000	362.210000	378.240000	352.890000

In [13]:

```
# Remove outliers from the data
# Calculate the IQR for each column
Q1 = data.quantile(0.25)
Q3 = data.quantile(0.75)
IQR = Q3 - Q1

# Define a threshold to identify outliers
outlier_threshold = 1.5

# Identify outliers using the IQR method
outliers = ((data < (Q1 - outlier_threshold * IQR)) | (data > (Q3 + outlier_threshold * IQR)))

# Remove outliers by replacing them with NaN
data_no_outliers = data[~outliers]

# Alternatively, you can choose to drop rows with any NaN values
data_no_outliers = data_no_outliers.dropna()

# Save the cleaned dataset
data_no_outliers.to_csv('cleaned_data.csv', index=False)
```

C:\Users\hp\AppData\Local\Temp\ipykernel\_15272\4198190489.py:3: FutureWarning: The default value of numeric\_only in DataFrame.quantile is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric\_only to silence this warning.

```
Q1 = data.quantile(0.25)
```

C:\Users\hp\AppData\Local\Temp\ipykernel\_15272\4198190489.py:4: FutureWarning: The default value of numeric\_only in DataFrame.quantile is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric\_only to silence this warning.

```
Q3 = data.quantile(0.75)
```

C:\Users\hp\AppData\Local\Temp\ipykernel\_15272\4198190489.py:11: FutureWarning: Automatic reindexing on DataFrame vs Series comparisons is deprecated and will raise ValueError in a future version. Do `left, right = left.align(right, axis=1, copy=False)` before e.g. `left == right`

```
outliers = ((data < (Q1 - outlier_threshold * IQR)) | (data > (Q3 + outlier_threshold * IQR)))
```

In [14]:

```
#Load the clean data  
data=pd.read_csv('cleaned_data.csv')  
#display the first ten entries  
data.head(10)
```

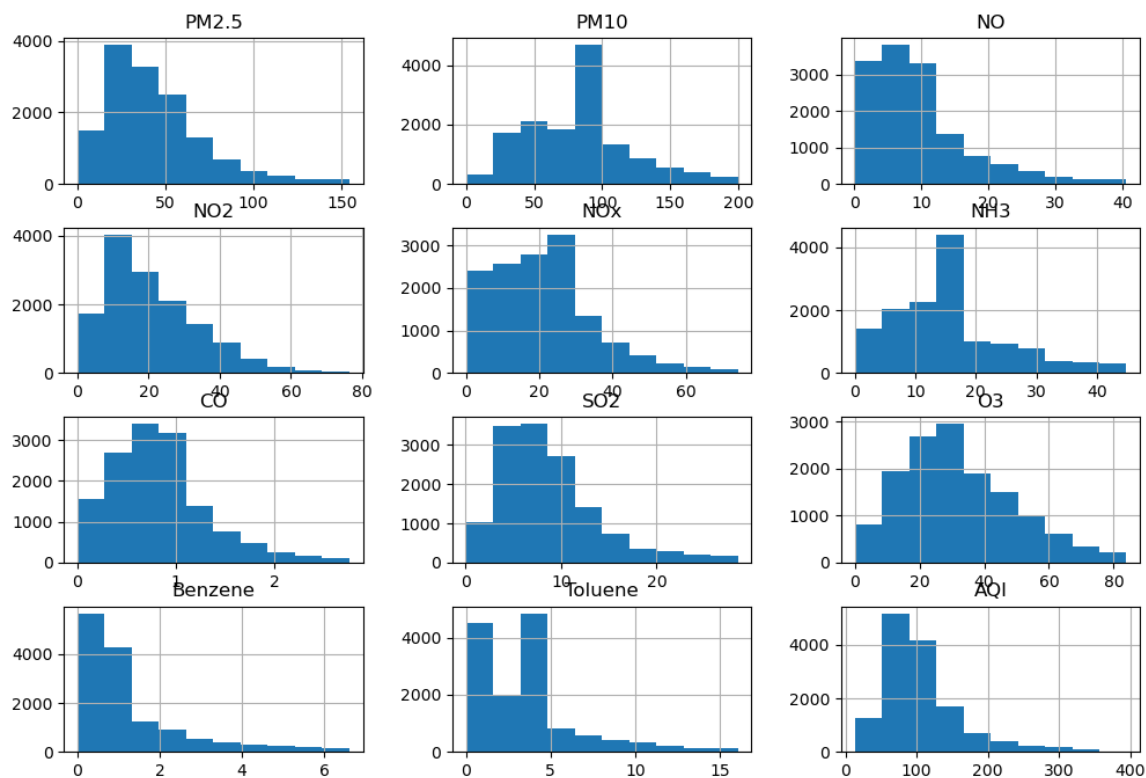
Out[14]:

	City	Date	PM2.5	PM10	NO	NO2	NOx	NH3	CO	SO2	O3	Benzene	T
0	Ahmedabad	2015-02-04	80.65	96.18	2.37	22.83	24.00	16.31	2.37	25.73	47.30	0.00	
1	Ahmedabad	2015-05-23	38.88	96.18	2.18	16.28	17.74	16.31	2.18	26.40	12.47	3.00	
2	Ahmedabad	2015-07-16	65.37	96.18	0.64	8.07	8.57	16.31	0.64	16.31	8.02	1.63	
3	Ahmedabad	2015-07-17	56.28	96.18	0.60	8.02	8.51	16.31	0.60	18.93	6.20	1.52	
4	Ahmedabad	2015-07-18	48.17	96.18	0.65	8.00	8.55	16.31	0.65	18.99	7.97	1.23	
5	Ahmedabad	2015-07-19	33.56	96.18	0.72	7.89	8.56	16.31	0.72	11.28	10.18	0.61	
6	Ahmedabad	2015-07-20	31.30	96.18	0.55	8.10	8.49	16.31	0.55	10.29	7.44	0.52	
7	Ahmedabad	2015-07-22	57.11	96.18	0.55	8.28	8.42	16.31	0.55	14.58	2.60	2.81	
8	Ahmedabad	2015-07-23	34.72	96.18	0.10	9.73	7.88	16.31	0.10	13.62	4.87	1.44	
9	Ahmedabad	2015-07-24	29.22	96.18	0.06	9.57	7.77	16.31	0.06	13.58	5.08	1.37	



In [15]:

```
# Visualize the data using histograms
import matplotlib.pyplot as plt
data.hist(figsize=(12, 8))
plt.show()
```





In [16]:

```
# Visualize the data using a correlation matrix heatmap
corr_matrix = data.corr(numeric_only=True)
fig=px.imshow(corr_matrix, title='Correlation Matrix')
fig.show()
```

In [17]:

```
#split the data into training and testing
X=data.drop(['City', 'Date', 'AQI', 'AQI_Bucket'],axis=1)
y=data['AQI']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=12)
```



In [18]:

```
#train the models
# Evaluation function
def evaluate_model(model, X_test, y_test):
    y_pred = model.predict(X_test)
    mae = mean_absolute_error(y_test, y_pred)
    mse = mean_squared_error(y_test, y_pred)
    rmse = np.sqrt(mse)
    r2 = r2_score(y_test, y_pred)

    return mae, mse, rmse, r2

# 1. Simple Linear Regression (using PM2.5 feature)
simple_linear_reg = LinearRegression()
simple_linear_reg.fit(X_train[['PM2.5']], y_train)
simple_linear_reg_metrics = evaluate_model(simple_linear_reg, X_test[['PM2.5']], y_test)

# 2. Multiple Linear Regression
multiple_linear_reg = LinearRegression()
multiple_linear_reg.fit(X_train, y_train)
multiple_linear_reg_metrics=evaluate_model(multiple_linear_reg,X_test,y_test)

# 3. Polynomial Regression
poly_features = PolynomialFeatures(degree=2)
X_train_poly = poly_features.fit_transform(X_train)
X_test_poly = poly_features.transform(X_test)
poly_reg = LinearRegression()
poly_reg.fit(X_train_poly, y_train)
poly_reg_metrics=evaluate_model(poly_reg,X_test_poly,y_test)

# 4. Ridge Regression
ridge_reg = Ridge(alpha=1.0)
ridge_reg.fit(X_train, y_train)
ridge_reg_metrics=evaluate_model(ridge_reg,X_test,y_test)

# 5. Lasso Regression
lasso_reg = Lasso(alpha=1.0)
lasso_reg.fit(X_train, y_train)
lasso_reg_metrics=evaluate_model(lasso_reg,X_test,y_test)

# 6. Support Vector Regression
svr_reg = SVR(kernel='linear')
svr_reg.fit(X_train, y_train)
svr_reg_metrics=evaluate_model(svr_reg,X_test,y_test)

# 7. Decision Tree Regression
dt_reg = DecisionTreeRegressor(random_state=12)
dt_reg.fit(X_train, y_train)
dt_reg_metrics=evaluate_model(dt_reg,X_test,y_test)

# 8. Random Forest Regression
rf_reg = RandomForestRegressor(random_state=12)
rf_reg.fit(X_train, y_train)
rf_reg_metrics=evaluate_model(rf_reg,X_test,y_test)

# Evaluate all models
models = ['Simple LR', 'Multiple LR', 'Polynomial LR', 'Ridge', 'Lasso', 'SVR', 'Decision
metrics = [simple_linear_reg_metrics, multiple_linear_reg_metrics, poly_reg_metrics, ridg
results = pd.DataFrame(metrics, columns=['MAE', 'MSE', 'RMSE', 'R2'], index=models)
```

```
print(results)
```

	MAE	MSE	RMSE	R2
Simple LR	20.943821	846.704924	29.098195	0.728816
Multiple LR	19.117010	735.675199	27.123333	0.764377
Polynomial LR	15.648078	584.202988	24.170291	0.812891
Ridge	19.116987	735.672527	27.123284	0.764378
Lasso	19.210513	740.873741	27.218996	0.762712
SVR	18.695228	776.086387	27.858327	0.751434
Decision Tree	20.987106	1139.842247	33.761550	0.634930
Random Forest	14.250897	543.056822	23.303580	0.826069

In [19]:

```
# Tune the best performing model in this case Random Forest regression

# Define the parameter grid for hyperparameter tuning
param_grid_rf = {
    'n_estimators': [100, 200, 300],
    'max_depth': [None, 10, 20],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}

# Create the Random Forest Regression model
rf_reg = RandomForestRegressor(random_state=12)

# Perform GridSearchCV for hyperparameter tuning
grid_search_rf = GridSearchCV(estimator=rf_reg, param_grid=param_grid_rf, cv=5, n_jobs=-1)
grid_search_rf.fit(X_train, y_train)

# Get the best model and its evaluation metrics on the test set
best_rf_reg = grid_search_rf.best_estimator_
best_rf_reg_metrics = evaluate_model(best_rf_reg, X_test, y_test)

print("Best hyperparameters for Random Forest:", grid_search_rf.best_params_)
print("Best score for Random Forest:", -grid_search_rf.best_score_)
```

```
Best hyperparameters for Random Forest: {'max_depth': None, 'min_samples_1
eaf': 4, 'min_samples_split': 2, 'n_estimators': 300}
Best score for Random Forest: -0.8215618059503266
```

In [ ]:

```

#Tune model using Randomized search
# Define the hyperparameter distribution
param_dist = {
    'n_estimators': randint(10, 200),
    'max_depth': randint(1, 30),
    'min_samples_split': randint(2, 10)
}

# Create a random forest regressor
rforest = RandomForestRegressor(random_state=12)

# Instantiate the randomized search with cross-validation
kfold = KFold(n_splits=5, shuffle=True, random_state=12)
random_search = RandomizedSearchCV(estimator=rforest, param_distributions=param_dist, n_i

# Fit the randomized search to the data
random_search.fit(X_train, y_train)

# Print the best combination of hyperparameters and their score
print("Best hyperparameters:", random_search.best_params_)
print("Best score:", -random_search.best_score_)

```

In [21]:

```


# Evaluate all models
models = ['Simple LR', 'Multiple LR', 'Polynomial LR', 'Ridge', 'Lasso', 'SVR', 'Decision
metrics = [simple_linear_reg_metrics, multiple_linear_reg_metrics, poly_reg_metrics, ridg
results = pd.DataFrame(metrics, columns=['MAE', 'MSE', 'RMSE', 'R2'], index=models)
print(results)

```

	MAE	MSE	RMSE	R2
Simple LR	20.943821	846.704924	29.098195	0.728816
Multiple LR	19.117010	735.675199	27.123333	0.764377
Polynomial LR	15.648078	584.202988	24.170291	0.812891
Ridge	19.116987	735.672527	27.123284	0.764378
Lasso	19.210513	740.873741	27.218996	0.762712
SVR	18.695228	776.086387	27.858327	0.751434
Decision Tree	20.987106	1139.842247	33.761550	0.634930
Random Forest	14.250897	543.056822	23.303580	0.826069
Tuned random forest	14.042782	537.855046	23.191702	0.827735

In [22]:

```
# Visualize the performance of each model using a bar plot
import matplotlib.pyplot as plt
results.plot(kind='bar', figsize=(12, 8), ylabel='Metric Value', title='Model Performance')
plt.show()
```

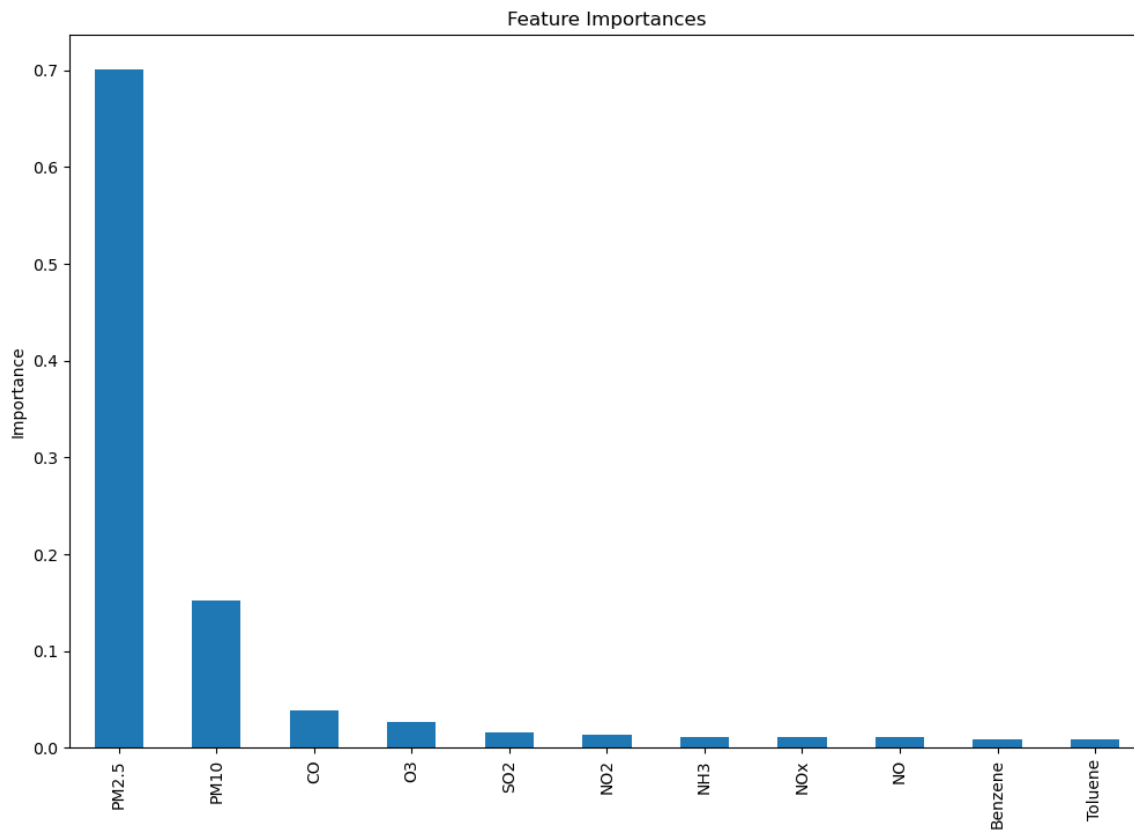


In [23]:

```
# Select the best model (e.g., tuned random forest)
best_model = best_rf_reg
```

In [25]:

```
# Analyze feature importances for the Random Forest model
importances = best_model.feature_importances_
feature_importances = pd.Series(importances, index=X.columns)
feature_importances.sort_values(ascending=False).plot(kind='bar', figsize=(12, 8), ylabel=
plt.show()
```



In [26]:

```
#save the best model
import joblib
joblib.dump(best_model, 'regression_model.pkl')
```

Out[26]:

```
['regression_model.pkl']
```

In [ ]:

In [ ]: